

```
hive> create external table ext_flight_1(fno int,src string,dest
string,start string,end2 string)
> stored by 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
> with serdeproperties("hbase.columns.mapping"=
":key,finfo:src,finfo:dest,ftime:start,ftime:end")
> tblproperties("hbase.table.name"='flight1');
```

OK

Time taken: 1.678 seconds

```
hive> select * from table_ext_flight_1
```

```
> ;
```

FAILED: SemanticException [Error 10001]: Line 1:14 Table not found

'table\_ext\_flight\_1'

```
hive> select * from ext_flight_1
```

```
> ;
```

OK

```
1  Pune  Mumbai 9.00  11.00
```

```
2  Delhi  Bangalore 10.00 12.00
```

Time taken: 0.493 seconds, Fetched: 2 row(s)

```
hive> help select
```

```
> ;
```

NoViableAltException(26@[])

at org.apache.hadoop.hive.ql.parse.HiveParser.statement(HiveParser.java:1020)

at org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:199)

at org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:166)

at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:393)

at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:307)

at org.apache.hadoop.hive.ql.Driver.compileInternal(Driver.java:1110)

at org.apache.hadoop.hive.ql.Driver.runInternal(Driver.java:1158)

at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1047)

```
at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1037)
at org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:207)
at org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:159)
at org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:370)
at org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:756)
at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:675)
at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:615)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:606)
at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
```

FAILED: ParseException line 1:0 cannot recognize input near 'help'

'select' '<EOF>'

hive> create table managed\_custInfo(id int,name string,orderid int)

row format delimited fields separated by "

;

>

>;

hive> create table managed\_custInfo(id int,name string,orderid int)

row format delimited fields terminated by ',' stored as textfile;

OK

Time taken: 0.099 seconds

hive> load data local inpath "/home/cloudera/custInfo.txt" into table

managed\_custInfo

>;

Loading data to table default.managed\_custinfo

Table default.managed\_custinfo stats: [numFiles=1, totalSize=24]

OK

Time taken: 0.904 seconds

```
hive> select * from managed_custInfo;
```

OK

```
1    "K"    2
```

```
3    "O"    4
```

```
5    "D"    6
```

Time taken: 0.096 seconds, Fetched: 3 row(s)

```
hive> select id from managed_custInfo;
```

OK

```
1
```

```
3
```

```
5
```

Time taken: 0.085 seconds, Fetched: 3 row(s)

```
hive> create table managed_orderInfo(orderid int,itemid int,quantity
int) row format delimited fields terminated by ',' stored as textfile;
```

OK

Time taken: 0.056 seconds

```
hive> load data local inpath "/home/cloudera/orderInfo.txt" into table
managed_orderInfo
```

```
>;
```

Loading data to table default.managed\_orderinfo

Table default.managed\_orderinfo stats: [numFiles=1, totalSize=18]

OK

Time taken: 0.566 seconds

```
hive> select * from managed_orderinfo;
```

OK

```
1    4    8
```

```
2    5    7
```

3 6 3

Time taken: 0.051 seconds, Fetched: 3 row(s)

```
hive> create table managed_itemInfo(itemid int,itemname
string,itemprice int) row format delimited fields terminated by ","
stored as textfile;
```

OK

Time taken: 0.072 seconds

```
hive> load data local inpath "/home/cloudera/itemInfo.txt" into table
managed_itemInfo
```

```
>;
```

Loading data to table default.managed\_iteminfo

Table default.managed\_iteminfo stats: [numFiles=1, totalSize=27]

OK

Time taken: 0.587 seconds

```
hive> select * from managed_itemInfo;
```

OK

```
4  "A"  10
```

```
5  "b"  20
```

```
1  "c"  30
```

Time taken: 0.062 seconds, Fetched: 3 row(s)

```
hive> select * from managed_customerInfo c innerjoin managed_orderInfo
o on (c.orderid=o.orderid);
```

FAILED: ParseException line 1:37 missing EOF at 'innerjoin' near 'c'

```
hive> select * from managed_customerInfo c join managed_orderInfo o on
(c.orderid=o.orderid);
```

FAILED: SemanticException [Error 10001]: Line 1:14 Table not found  
'managed\_customerInfo'

```
hive> select * from managed_custInfo c join managed_orderInfo o on
(c.orderid=o.orderid);
```

Query ID = cloudera\_20220515120505\_f9306380-934d-4858-a011-484792dec2c6

Total jobs = 1

Execution log at:

/tmp/cloudera/cloudera\_20220515120505\_f9306380-934d-4858-a011-484792dec2c6.log

2022-05-15 12:05:10 Starting to launch local task to process map

join; maximum memory = 1013645312

2022-05-15 12:05:11 Dump the side-table for tag: 1 with group count: 3

into file: file:/tmp/cloudera/91eaf19a-4a2c-4dba-b6cc-25f8ba771942/hive\_2022-05-15\_12-05-06\_582\_7133824769850861002-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile01--.hashtable

2022-05-15 12:05:11 Uploaded 1 File to:

file:/tmp/cloudera/91eaf19a-4a2c-4dba-b6cc-25f8ba771942/hive\_2022-05-15\_12-05-06\_582\_7133824769850861002-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile01--.hashtable  
(323 bytes)

2022-05-15 12:05:11 End of local task; Time Taken: 1.101 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_1652600087425\_0001, Tracking URL =

[http://quickstart.cloudera:8088/proxy/application\\_1652600087425\\_0001/](http://quickstart.cloudera:8088/proxy/application_1652600087425_0001/)

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1652600087425\_0001

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2022-05-15 12:05:21,644 Stage-3 map = 0%, reduce = 0%

2022-05-15 12:05:28,177 Stage-3 map = 100%, reduce = 0%, Cumulative

CPU 1.01 sec

MapReduce Total cumulative CPU time: 1 seconds 10 msec

Ended Job = job\_1652600087425\_0001

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 1.01 sec HDFS Read: 6478

HDFS Write: 14 SUCCESS

Total MapReduce CPU Time Spent: 1 seconds 10 msec

OK

1    "K"    2    2    5    7

Time taken: 22.69 seconds, Fetched: 1 row(s)

hive> select \* from managed\_custInfo c join managed\_orderInfo o on

(c.orderid=o.orderid);

Query ID = cloudera\_20220515121616\_3c8272db-a221-460e-9d69-c0fc116268d0

Total jobs = 1

Interrupting... Be patient, this might take some time.

Press Ctrl+C again to kill JVM

Execution failed with exit status: 130

Obtaining error information

Task failed!

Task ID:

Stage-4

Logs:

/tmp/cloudera/hive.log

FAILED: Execution Error, return code 130 from

org.apache.hadoop.hive ql.exec.mr.MapredLocalTask

hive> select o.orderid,i.itemid,i.itemname,i.itemprice\*o.quantity as

totalPrice from managed\_itemInfo i join managed\_orderInfo o on

(i.itemid=o.itemid);

Query ID = cloudera\_20220515121919\_73b3661b-bd6b-4a55-8359-cb4356bd19f3

Total jobs = 1

Execution log at:

/tmp/cloudera/cloudera\_20220515121919\_73b3661b-bd6b-4a55-8359-cb4356bd19f3.log

2022-05-15 12:19:16 Starting to launch local task to process map

join; maximum memory = 1013645312

2022-05-15 12:19:18 Dump the side-table for tag: 1 with group count: 3

into file: file:/tmp/cloudera/91eaf19a-4a2c-4dba-b6cc-25f8ba771942/hive\_2022-05-15\_12-19-11\_869\_2957621910429033491-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile21--.hashtable

2022-05-15 12:19:18 Uploaded 1 File to:

file:/tmp/cloudera/91eaf19a-4a2c-4dba-b6cc-25f8ba771942/hive\_2022-05-15\_12-19-11\_869\_2957621910429033491-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile21--.hashtable  
(323 bytes)

2022-05-15 12:19:18 End of local task; Time Taken: 1.971 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_1652600087425\_0002, Tracking URL =

[http://quickstart.cloudera:8088/proxy/application\\_1652600087425\\_0002/](http://quickstart.cloudera:8088/proxy/application_1652600087425_0002/)

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1652600087425\_0002

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2022-05-15 12:19:29,119 Stage-3 map = 0%, reduce = 0%

2022-05-15 12:19:38,736 Stage-3 map = 100%, reduce = 0%, Cumulative

CPU 1.72 sec

MapReduce Total cumulative CPU time: 1 seconds 720 msec

Ended Job = job\_1652600087425\_0002

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 1.72 sec HDFS Read: 6933

HDFS Write: 23 SUCCESS

Total MapReduce CPU Time Spent: 1 seconds 720 msec

OK

```
1  4  "A"  80
2  5  "b" 140
```

Time taken: 27.942 seconds, Fetched: 2 row(s)

```
hive> select quantity,orderid from orderInfo;
```

FAILED: SemanticException [Error 10001]: Line 1:29 Table not found 'orderInfo'

```
hive> select quantity,orderid from managed_orderInfo;
```

OK

```
8  1
7  2
3  3
```

Time taken: 0.062 seconds, Fetched: 3 row(s)

```
hive> select o.orderid,i.itemid,i.itemname,i.itemprice*o.quantity as
totalPrice,max(totalPrice) from managed_itemInfo i join
managed_orderInfo o on (i.itemid=o.itemid);
```

FAILED: SemanticException [Error 10004]: Line 1:78 Invalid table alias  
or column reference 'totalPrice': (possible column names are:  
i.itemid, i.itemname, i.itemprice, o.orderid, o.itemid, o.quantity)

```
hive> select o.orderid,i.itemid,i.itemname,max(i.itemprice*o.quantity)
as maxtotalPrice from managed_itemInfo i join managed_orderInfo o on
(i.itemid=o.itemid);
```

FAILED: SemanticException [Error 10025]: Line 1:7 Expression not in  
GROUP BY key 'orderid'

```
hive> select o.orderid,i.itemid,i.itemname,i.itemprice*o.quantity as
max(totalPrice) from managed_itemInfo i join managed_orderInfo o on
(i.itemid=o.itemid);
```

FAILED: ParseException line 1:66 missing EOF at '(' near 'max'

```
hive> select o.orderid,i.itemid,i.itemname,max(i.itemprice*o.quantity)
from managed_itemInfo i join managed_orderInfo o on
(i.itemid=o.itemid);
```



FAILED: SemanticException [Error 10025]: Line 1:7 Expression not in  
GROUP BY key 'orderid'

```
hive> select max(i.itemprice*o.quantity) from managed_itemInfo i join  
managed_orderInfo o on (i.itemid=o.itemid);
```

Query ID = cloudera\_20220515122828\_8639b14c-5f7a-425c-b11c-24fe5f272ddb

Total jobs = 1

Execution log at:

/tmp/cloudera/cloudera\_20220515122828\_8639b14c-5f7a-425c-b11c-24fe5f272ddb.log

2022-05-15 12:28:37 Starting to launch local task to process map

join; maximum memory = 1013645312

2022-05-15 12:28:39 Dump the side-table for tag: 1 with group count: 3

into file: file:/tmp/cloudera/91eaf19a-4a2c-4dba-b6cc-25f8ba771942/hive\_2022-05-15\_12-28-  
33\_236\_779602107621161755-1/-local-10004/HashTable-Stage-2/MapJoin-mapfile31--.hashtable

2022-05-15 12:28:39 Uploaded 1 File to:

file:/tmp/cloudera/91eaf19a-4a2c-4dba-b6cc-25f8ba771942/hive\_2022-05-15\_12-28-  
33\_236\_779602107621161755-1/-local-10004/HashTable-Stage-2/MapJoin-mapfile31--.hashtable  
(320 bytes)

2022-05-15 12:28:39 End of local task; Time Taken: 1.845 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1652600087425\_0003, Tracking URL =

[http://quickstart.cloudera:8088/proxy/application\\_1652600087425\\_0003/](http://quickstart.cloudera:8088/proxy/application_1652600087425_0003/)

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1652600087425\_0003

Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1

2022-05-15 12:28:47,692 Stage-2 map = 0%, reduce = 0%

2022-05-15 12:28:56,256 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 1.4 sec

2022-05-15 12:29:04,703 Stage-2 map = 100%, reduce = 100%, Cumulative

CPU 2.66 sec

MapReduce Total cumulative CPU time: 2 seconds 660 msec

Ended Job = job\_1652600087425\_0003

MapReduce Jobs Launched:

Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 2.66 sec HDFS

Read: 11363 HDFS Write: 4 SUCCESS

Total MapReduce CPU Time Spent: 2 seconds 660 msec

OK

140

Time taken: 32.555 seconds, Fetched: 1 row(s)

hive>